

learning-using-keras-regression

November 21, 2024

Importing the necessary libraries and packages

```
[12]: import warnings
      warnings.filterwarnings('ignore')
      from sklearn.datasets import fetch_california_housing
      from sklearn.model_selection import train_test_split
      import tensorflow as tf
```

```
[31]: print(tf.__version__)
```

2.18.0

Downloading the dataset

- California Housing dataset

```
[2]: housing = fetch_california_housing()
```

Inspecting the data and it's shape

```
[3]: housing.data, housing.data.shape
```

```
[3]: (array([[ 8.3252    ,  41.        ,  6.98412698, ...,  2.55555556,
            37.88    , -122.23    ],
           [ 8.3014    ,  21.        ,  6.23813708, ...,  2.10984183,
            37.86    , -122.22    ],
           [ 7.2574    ,  52.        ,  8.28813559, ...,  2.80225989,
            37.85    , -122.24    ],
           ...,
           [ 1.7       ,  17.        ,  5.20554273, ...,  2.3256351 ,
            39.43    , -121.22    ],
           [ 1.8672    ,  18.        ,  5.32951289, ...,  2.12320917,
            39.43    , -121.32    ],
           [ 2.3886    ,  16.        ,  5.25471698, ...,  2.61698113,
            39.37    , -121.24    ]]),
      (20640, 8))
```

Inspecting the data target and it's shape

```
[4]: housing.target, housing.target.shape
```

```
[4]: (array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]), (20640,))
```

Splitting the dataset

- Splitting the dataset into training and testing dataset (75:25 percentage split)

```
[5]: X_train_full, X_test, y_train_full, y_test = train_test_split(housing.data, ↵  
    ↪housing.target, random_state=42)
```

```
[6]: print(X_train_full.shape, y_train_full.shape)  
     print(X_test.shape, y_test.shape)
```

```
(15480, 8) (15480,)
```

```
(5160, 8) (5160,)
```

- Splitting the training dataset into training and validation dataset (75:25 percentage split)

```
[7]: X_train, X_valid, y_train, y_valid = train_test_split(X_train_full, ↵  
    ↪y_train_full, random_state=42)
```

```
[8]: print(X_train.shape, y_train.shape)  
     print(X_valid.shape, y_valid.shape)
```

```
(11610, 8) (11610,)
```

```
(3870, 8) (3870,)
```

0.0.1 Type 1 Model : Initializing the model

- Flatten() layer is not used in this example.
- Instead a Normalization layer is used as the first layer and does the same thing as Scikit-Learn's StandardScaler().
- Must be fitted to the training data using its adapt() methods before calling the model's fit() method

The **Normalization layer** learns the feature means and standard deviations in the training data when you call the adapt() method. Yet when you display the model's summary, these statistics are listed as non trainable. This is because these parameters are not affected by gradient descent.

```
[9]: tf.random.set_seed(42)
```

```
[21]: norm_layer = tf.keras.layers.Normalization(input_shape=X_train.shape[1:])  
     model = tf.keras.Sequential([  
         norm_layer,  
         tf.keras.layers.Dense(50, activation='relu'),  
         tf.keras.layers.Dense(50, activation='relu'),  
         tf.keras.layers.Dense(50, activation='relu'),  
         tf.keras.layers.Dense(1)])
```

Setting the optimizer

```
[22]: optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)
```

Setting the model's training metrics

```
[23]: model.compile(loss="mse", optimizer=optimizer, metrics=["RootMeanSquaredError"])
```

Adapt the training data to the normalization layer

```
[24]: norm_layer.adapt(X_train)
```

Train the model on training data using specified number of epochs

```
[25]: history = model.fit(X_train, y_train, epochs=20, validation_data=(X_valid,   
    ↪ y_valid))
```

```
Epoch 1/20
363/363          1s 1ms/step -
RootMeanSquaredError: 1.1591 - loss: 1.4237 - val_RootMeanSquaredError: 0.9103 -
val_loss: 0.8287
Epoch 2/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.6435 - loss: 0.4146 - val_RootMeanSquaredError: 0.6084 -
val_loss: 0.3701
Epoch 3/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.6060 - loss: 0.3674 - val_RootMeanSquaredError: 0.8837 -
val_loss: 0.7809
Epoch 4/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5923 - loss: 0.3510 - val_RootMeanSquaredError: 1.1057 -
val_loss: 1.2226
Epoch 5/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5820 - loss: 0.3389 - val_RootMeanSquaredError: 1.0742 -
val_loss: 1.1540
Epoch 6/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5759 - loss: 0.3317 - val_RootMeanSquaredError: 1.5071 -
val_loss: 2.2712
Epoch 7/20
363/363          0s 934us/step -
RootMeanSquaredError: 0.5718 - loss: 0.3271 - val_RootMeanSquaredError: 0.9234 -
val_loss: 0.8526
Epoch 8/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5648 - loss: 0.3191 - val_RootMeanSquaredError: 0.8959 -
val_loss: 0.8026
Epoch 9/20
```

```

363/363          0s 988us/step -
RootMeanSquaredError: 0.5563 - loss: 0.3096 - val_RootMeanSquaredError: 0.5539 -
val_loss: 0.3068
Epoch 10/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5518 - loss: 0.3046 - val_RootMeanSquaredError: 0.5721 -
val_loss: 0.3273
Epoch 11/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5476 - loss: 0.3000 - val_RootMeanSquaredError: 0.5610 -
val_loss: 0.3147
Epoch 12/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5432 - loss: 0.2952 - val_RootMeanSquaredError: 0.5622 -
val_loss: 0.3161
Epoch 13/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5392 - loss: 0.2909 - val_RootMeanSquaredError: 0.5681 -
val_loss: 0.3227
Epoch 14/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5364 - loss: 0.2878 - val_RootMeanSquaredError: 0.5309 -
val_loss: 0.2819
Epoch 15/20
363/363          0s 976us/step -
RootMeanSquaredError: 0.5329 - loss: 0.2841 - val_RootMeanSquaredError: 0.5499 -
val_loss: 0.3023
Epoch 16/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5316 - loss: 0.2827 - val_RootMeanSquaredError: 0.5311 -
val_loss: 0.2821
Epoch 17/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5323 - loss: 0.2834 - val_RootMeanSquaredError: 0.6966 -
val_loss: 0.4852
Epoch 18/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5275 - loss: 0.2783 - val_RootMeanSquaredError: 0.5416 -
val_loss: 0.2933
Epoch 19/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5250 - loss: 0.2757 - val_RootMeanSquaredError: 0.5675 -
val_loss: 0.3221
Epoch 20/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5224 - loss: 0.2730 - val_RootMeanSquaredError: 0.5398 -
val_loss: 0.2914

```

Evaluate the model's performance on testing data

```
[26]: mse_test, rmse_test = model.evaluate(X_test, y_test)
```

```
162/162          0s 652us/step -  
RootMeanSquaredError: 0.5293 - loss: 0.2803
```

```
[27]: print(mse_test, rmse_test)
```

```
0.2853373885154724 0.5341697931289673
```

Predict the output values for the first 3 datapoints

```
[28]: X_new = X_test[:3]  
y_pred = model.predict(X_new)
```

```
1/1          0s 60ms/step
```

```
[29]: y_pred
```

```
[29]: array([[0.45630956],  
            [1.4869951 ],  
            [5.008497  ]], dtype=float32)
```

0.0.2 Type 2 Model : Building a Wide and Deep neural network

- Non sequential neural network is a Wide and Deep Neural Network
- Heng-Tze Cheng et al (2016 paper)
- It connects all or part of the inputs directly to the output layer.
- This architecture makes it possible for the neural network to learn both deep patterns (using the deep path) and simple rules(through the short path)
- A regular MLP forces all the data to flow through the full stack of layers, thus simple patterns in the data may end up being distorted by this sequence of transformations

Defining layers of the model

1. Normalization layer to standardize the inputs
2. Two Dense layers with 30 neurons each using the ReLU activation function
3. Concatenate layer
4. Dense layer with a single neuron for the output layer, without any activation function

```
[33]: normalization_layer = tf.keras.layers.Normalization()  
hidden_layer1 = tf.keras.layers.Dense(30, activation='relu')  
hidden_layer2 = tf.keras.layers.Dense(30, activation='relu')  
concat_layer = tf.keras.layers.Concatenate()  
output_layer = tf.keras.layers.Dense(1)
```

Establishing the flow of data in model layers

1. `input_` - Input object. It is a specification of the kind of input the model will get, including its shape and optionally its dtype, which defaults to 32 bits floats. A model may actually have multiple inputs. Input object is just a data specification.
2. Used the Normalization layer just like a function, passing it the Input object. It's called the functional API. No actual data is being processed yet. Only keras is being told how it should connect the layers together. The input and output are both symbolic. Normalized doesn't store any actual data, it's just used to construct the model.
3. `concat_layer` - concatenate the input and the second hidden layer's output.

```
[34]: input_ = tf.keras.layers.Input(shape=X_train.shape[1:])
normalized = normalization_layer(input_)
hidden1 = hidden_layer1(normalized)
hidden2 = hidden_layer2(hidden1)
concat = concat_layer([normalized, hidden2])
output = output_layer(concat)
```

Creation of Keras model from inputs and outputs

```
[38]: model = tf.keras.Model(inputs=[input_], outputs=[output])
```

Defining the optimizer

```
[39]: optimizer1 = tf.keras.optimizers.Adam(learning_rate=1e-3)
```

Model compilation using model training metrics

```
[40]: model.compile(loss="mse", optimizer=optimizer1,
↳ metrics=["RootMeanSquaredError"])
```

Fitting the model

```
[41]: history = model.fit(X_train, y_train, epochs=20, validation_data=(X_valid,
↳ y_valid))
```

```
Epoch 1/20
363/363          1s 1ms/step -
RootMeanSquaredError: 98.5840 - loss: 13144.9893 - val_RootMeanSquaredError:
13.7108 - val_loss: 187.9848
Epoch 2/20
363/363          0s 1ms/step -
RootMeanSquaredError: 2.6266 - loss: 6.9281 - val_RootMeanSquaredError: 12.9485
- val_loss: 167.6648
Epoch 3/20
363/363          0s 950us/step -
RootMeanSquaredError: 1.9082 - loss: 3.6802 - val_RootMeanSquaredError: 11.5206
- val_loss: 132.7237
Epoch 4/20
363/363          0s 882us/step -
RootMeanSquaredError: 1.5835 - loss: 2.5444 - val_RootMeanSquaredError: 10.2225
```

```

- val_loss: 104.4997
Epoch 5/20
363/363          0s 966us/step -
RootMeanSquaredError: 1.3602 - loss: 1.8849 - val_RootMeanSquaredError: 8.9722 -
val_loss: 80.5004
Epoch 6/20
363/363          0s 1ms/step -
RootMeanSquaredError: 1.2263 - loss: 1.5357 - val_RootMeanSquaredError: 7.7842 -
val_loss: 60.5934
Epoch 7/20
363/363          0s 988us/step -
RootMeanSquaredError: 1.1364 - loss: 1.3156 - val_RootMeanSquaredError: 6.6950 -
val_loss: 44.8229
Epoch 8/20
363/363          0s 988us/step -
RootMeanSquaredError: 1.0486 - loss: 1.1081 - val_RootMeanSquaredError: 5.6728 -
val_loss: 32.1808
Epoch 9/20
363/363          0s 998us/step -
RootMeanSquaredError: 1.0188 - loss: 1.0466 - val_RootMeanSquaredError: 4.7339 -
val_loss: 22.4094
Epoch 10/20
363/363          0s 969us/step -
RootMeanSquaredError: 0.9584 - loss: 0.9214 - val_RootMeanSquaredError: 3.9133 -
val_loss: 15.3142
Epoch 11/20
363/363          0s 989us/step -
RootMeanSquaredError: 0.9036 - loss: 0.8168 - val_RootMeanSquaredError: 3.2180 -
val_loss: 10.3554
Epoch 12/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.8940 - loss: 0.7995 - val_RootMeanSquaredError: 2.6714 -
val_loss: 7.1362
Epoch 13/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.9756 - loss: 0.9531 - val_RootMeanSquaredError: 2.2724 -
val_loss: 5.1636
Epoch 14/20
363/363          0s 988us/step -
RootMeanSquaredError: 0.9971 - loss: 0.9992 - val_RootMeanSquaredError: 2.0833 -
val_loss: 4.3401
Epoch 15/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.9919 - loss: 0.9905 - val_RootMeanSquaredError: 2.2558 -
val_loss: 5.0888
Epoch 16/20
363/363          0s 977us/step -
RootMeanSquaredError: 1.0314 - loss: 1.0712 - val_RootMeanSquaredError: 1.9443 -

```

```

val_loss: 3.7805
Epoch 17/20
363/363          0s 945us/step -
RootMeanSquaredError: 0.9496 - loss: 0.9035 - val_RootMeanSquaredError: 2.4768 -
val_loss: 6.1343
Epoch 18/20
363/363          0s 979us/step -
RootMeanSquaredError: 1.3026 - loss: 1.7384 - val_RootMeanSquaredError: 2.3679 -
val_loss: 5.6067
Epoch 19/20
363/363          0s 965us/step -
RootMeanSquaredError: 1.4613 - loss: 2.2241 - val_RootMeanSquaredError: 2.0035 -
val_loss: 4.0142
Epoch 20/20
363/363          0s 1ms/step -
RootMeanSquaredError: 1.0357 - loss: 1.0998 - val_RootMeanSquaredError: 2.1053 -
val_loss: 4.4321

```

Evaluate the model's performance on testing data

```
[42]: mse_test, rmse_test = model.evaluate(X_test, y_test)
```

```

162/162          0s 845us/step -
RootMeanSquaredError: 1.3568 - loss: 1.8413

```

```
[43]: print(mse_test, rmse_test)
```

```
1.860201120376587 1.3638919591903687
```

0.0.3 Type 3 Model : Subset features different paths

- Send a subset of the features through the wide path and a different subset (possibly overlapping) through the deep path

Defining layers of the model

```
[52]: input_wide = tf.keras.layers.Input(shape=[5]) # features 0 to 4
input_deep = tf.keras.layers.Input(shape=[6]) # features 2 to 7

norm_layer_wide = tf.keras.layers.Normalization()
norm_layer_deep = tf.keras.layers.Normalization()

norm_wide = norm_layer_wide(input_wide)
norm_deep = norm_layer_deep(input_deep)

hidden1 = tf.keras.layers.Dense(30, activation='relu')(norm_deep)
hidden2 = tf.keras.layers.Dense(30, activation='relu')(hidden1)

concat = tf.keras.layers.concatenate([norm_wide, hidden2])

```



```
output = tf.keras.layers.Dense(1)(concat)

model = tf.keras.Model(inputs=[input_wide, input_deep], outputs=[output])
```

Defining the optimizer

```
[53]: optimizer2 = tf.keras.optimizers.Adam(learning_rate=1e-3)
```

Compiling the model with training metrics

```
[54]: model.compile(loss="mse", optimizer=optimizer2,
    ↪metrics=["RootMeanSquaredError"])
```

Seperating the datasets

```
[58]: X_train_wide, X_train_deep = X_train[:, :5], X_train[:, 2:]
X_valid_wide, X_valid_deep = X_valid[:, :5], X_valid[:, 2:]
X_test_wide, X_test_deep = X_test[:, :5], X_test[:, 2:]
X_new_wide, X_new_deep = X_test_wide[:3], X_test_deep[:3]
```

Training the model

```
[56]: norm_layer_wide.adapt(X_train_wide)
norm_layer_deep.adapt(X_train_deep)

history = model.fit((X_train_wide, X_train_deep), y_train, epochs=20,
    ↪validation_data=((X_valid_wide, X_valid_deep), y_valid))
```

Epoch 1/20

363/363 1s 1ms/step -

RootMeanSquaredError: 1.4704 - loss: 2.2448 - val_RootMeanSquaredError: 1.4442 -
val_loss: 2.0857

Epoch 2/20

363/363 0s 991us/step -

RootMeanSquaredError: 0.7428 - loss: 0.5526 - val_RootMeanSquaredError: 0.8360 -
val_loss: 0.6988

Epoch 3/20

363/363 0s 905us/step -

RootMeanSquaredError: 0.6676 - loss: 0.4461 - val_RootMeanSquaredError: 0.6239 -
val_loss: 0.3893

Epoch 4/20

363/363 0s 958us/step -

RootMeanSquaredError: 0.6452 - loss: 0.4165 - val_RootMeanSquaredError: 0.6018 -
val_loss: 0.3621

Epoch 5/20

363/363 0s 1ms/step -

RootMeanSquaredError: 0.6309 - loss: 0.3983 - val_RootMeanSquaredError: 0.5956 -
val_loss: 0.3547

Epoch 6/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.6214 - loss: 0.3864 - val_RootMeanSquaredError: 0.6294 -
val_loss: 0.3961
Epoch 7/20
363/363 0s 999us/step -
RootMeanSquaredError: 0.6138 - loss: 0.3770 - val_RootMeanSquaredError: 0.5915 -
val_loss: 0.3499
Epoch 8/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.6070 - loss: 0.3687 - val_RootMeanSquaredError: 0.6080 -
val_loss: 0.3697
Epoch 9/20
363/363 0s 996us/step -
RootMeanSquaredError: 0.6021 - loss: 0.3628 - val_RootMeanSquaredError: 0.5725 -
val_loss: 0.3278
Epoch 10/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5970 - loss: 0.3566 - val_RootMeanSquaredError: 0.6704 -
val_loss: 0.4494
Epoch 11/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5933 - loss: 0.3521 - val_RootMeanSquaredError: 0.6333 -
val_loss: 0.4011
Epoch 12/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5908 - loss: 0.3491 - val_RootMeanSquaredError: 0.8573 -
val_loss: 0.7349
Epoch 13/20
363/363 0s 955us/step -
RootMeanSquaredError: 0.5876 - loss: 0.3454 - val_RootMeanSquaredError: 0.7427 -
val_loss: 0.5515
Epoch 14/20
363/363 0s 915us/step -
RootMeanSquaredError: 0.5855 - loss: 0.3429 - val_RootMeanSquaredError: 1.1110 -
val_loss: 1.2343
Epoch 15/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5862 - loss: 0.3438 - val_RootMeanSquaredError: 0.8765 -
val_loss: 0.7682
Epoch 16/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5816 - loss: 0.3384 - val_RootMeanSquaredError: 1.2828 -
val_loss: 1.6457
Epoch 17/20
363/363 0s 1ms/step -
RootMeanSquaredError: 0.5817 - loss: 0.3385 - val_RootMeanSquaredError: 1.1154 -
val_loss: 1.2440

```
Epoch 18/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5795 - loss: 0.3359 - val_RootMeanSquaredError: 1.0683 -
val_loss: 1.1412
Epoch 19/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5771 - loss: 0.3332 - val_RootMeanSquaredError: 0.7588 -
val_loss: 0.5758
Epoch 20/20
363/363          0s 1ms/step -
RootMeanSquaredError: 0.5730 - loss: 0.3285 - val_RootMeanSquaredError: 0.7698 -
val_loss: 0.5926
```

Evaluating the model against test dataset

```
[59]: mse_test = model.evaluate((X_test_wide, X_test_deep), y_test)
      y_pred = model.predict((X_new_wide, X_new_deep))
```

```
162/162          0s 574us/step -
RootMeanSquaredError: 0.5748 - loss: 0.3305
1/1              0s 44ms/step
```

0.0.4 Type 4 Model : Seperate Outputs

- Extra output is quite easy: connect it to the appropriate layer and add it to the model's list of outputs

Defining the 2 model outputs in output layer

```
[71]: output = tf.keras.layers.Dense(1)(concat)
      aux_output = tf.keras.layers.Dense(1)(hidden2)

      model = tf.keras.Model(inputs=[input_wide, input_deep], outputs=[output,
      ↪aux_output])
```

Defining the optimizer

```
[72]: optimizer3 = tf.keras.optimizers.Adam(learning_rate=1e-3)
```

Define the model compilation with training parameters

```
[73]: model.compile(loss=("mse", "mse"), loss_weights=(0.9, 0.1), optimizer =
      ↪optimizer3, metrics=["RootMeanSquaredError", "RootMeanSquaredError"])
```

Adapting the training data with normalization layer

```
[74]: norm_layer_wide.adapt(X_train_wide)
      norm_layer_deep.adapt(X_train_deep)
```

Fitting the model

```
[75]: history = model.fit((X_train_wide, X_train_deep), (y_train, y_train),  
    ↪ epochs=20, validation_data=((X_valid_wide, X_valid_deep), (y_valid,  
    ↪ y_valid)))
```

Epoch 1/20

363/363 2s 2ms/step -

dense_29_RootMeanSquaredError: 1.4116 - dense_29_loss: 1.8856 -
dense_30_RootMeanSquaredError: 1.8370 - dense_30_loss: 0.3484 - loss: 2.2340 -
val_dense_29_RootMeanSquaredError: 0.7331 - val_dense_29_loss: 0.4836 -
val_dense_30_RootMeanSquaredError: 1.0084 - val_dense_30_loss: 0.1017 -
val_loss: 0.5854

Epoch 2/20

363/363 0s 1ms/step -

dense_29_RootMeanSquaredError: 0.7011 - dense_29_loss: 0.4427 -
dense_30_RootMeanSquaredError: 0.8633 - dense_30_loss: 0.0746 - loss: 0.5173 -
val_dense_29_RootMeanSquaredError: 0.6377 - val_dense_29_loss: 0.3659 -
val_dense_30_RootMeanSquaredError: 0.8684 - val_dense_30_loss: 0.0754 -
val_loss: 0.4414

Epoch 3/20

363/363 1s 1ms/step -

dense_29_RootMeanSquaredError: 0.6355 - dense_29_loss: 0.3636 -
dense_30_RootMeanSquaredError: 0.7202 - dense_30_loss: 0.0519 - loss: 0.4155 -
val_dense_29_RootMeanSquaredError: 0.6087 - val_dense_29_loss: 0.3334 -
val_dense_30_RootMeanSquaredError: 0.7227 - val_dense_30_loss: 0.0522 -
val_loss: 0.3857

Epoch 4/20

363/363 1s 1ms/step -

dense_29_RootMeanSquaredError: 0.6065 - dense_29_loss: 0.3311 -
dense_30_RootMeanSquaredError: 0.6699 - dense_30_loss: 0.0449 - loss: 0.3760 -
val_dense_29_RootMeanSquaredError: 0.5841 - val_dense_29_loss: 0.3071 -
val_dense_30_RootMeanSquaredError: 0.6526 - val_dense_30_loss: 0.0426 -
val_loss: 0.3497

Epoch 5/20

363/363 1s 1ms/step -

dense_29_RootMeanSquaredError: 0.5936 - dense_29_loss: 0.3172 -
dense_30_RootMeanSquaredError: 0.6485 - dense_30_loss: 0.0421 - loss: 0.3592 -
val_dense_29_RootMeanSquaredError: 0.5906 - val_dense_29_loss: 0.3139 -
val_dense_30_RootMeanSquaredError: 0.6445 - val_dense_30_loss: 0.0415 -
val_loss: 0.3555

Epoch 6/20

363/363 1s 1ms/step -

dense_29_RootMeanSquaredError: 0.5862 - dense_29_loss: 0.3093 -
dense_30_RootMeanSquaredError: 0.6372 - dense_30_loss: 0.0406 - loss: 0.3499 -
val_dense_29_RootMeanSquaredError: 0.6883 - val_dense_29_loss: 0.4263 -
val_dense_30_RootMeanSquaredError: 0.8057 - val_dense_30_loss: 0.0649 -
val_loss: 0.4913

Epoch 7/20
 363/363 1s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5817 - dense_29_loss: 0.3046 -
 dense_30_RootMeanSquaredError: 0.6302 - dense_30_loss: 0.0397 - loss: 0.3443 -
 val_dense_29_RootMeanSquaredError: 0.6775 - val_dense_29_loss: 0.4130 -
 val_dense_30_RootMeanSquaredError: 0.6516 - val_dense_30_loss: 0.0424 -
 val_loss: 0.4556
 Epoch 8/20
 363/363 0s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5791 - dense_29_loss: 0.3018 -
 dense_30_RootMeanSquaredError: 0.6240 - dense_30_loss: 0.0390 - loss: 0.3408 -
 val_dense_29_RootMeanSquaredError: 0.8057 - val_dense_29_loss: 0.5841 -
 val_dense_30_RootMeanSquaredError: 0.8235 - val_dense_30_loss: 0.0678 -
 val_loss: 0.6521
 Epoch 9/20
 363/363 1s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5762 - dense_29_loss: 0.2989 -
 dense_30_RootMeanSquaredError: 0.6193 - dense_30_loss: 0.0384 - loss: 0.3372 -
 val_dense_29_RootMeanSquaredError: 0.6289 - val_dense_29_loss: 0.3559 -
 val_dense_30_RootMeanSquaredError: 0.6241 - val_dense_30_loss: 0.0389 -
 val_loss: 0.3949
 Epoch 10/20
 363/363 1s 2ms/step -
 dense_29_RootMeanSquaredError: 0.5744 - dense_29_loss: 0.2970 -
 dense_30_RootMeanSquaredError: 0.6154 - dense_30_loss: 0.0379 - loss: 0.3348 -
 val_dense_29_RootMeanSquaredError: 0.7995 - val_dense_29_loss: 0.5751 -
 val_dense_30_RootMeanSquaredError: 0.8740 - val_dense_30_loss: 0.0764 -
 val_loss: 0.6517
 Epoch 11/20
 363/363 1s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5722 - dense_29_loss: 0.2947 -
 dense_30_RootMeanSquaredError: 0.6129 - dense_30_loss: 0.0376 - loss: 0.3323 -
 val_dense_29_RootMeanSquaredError: 0.5750 - val_dense_29_loss: 0.2975 -
 val_dense_30_RootMeanSquaredError: 0.6114 - val_dense_30_loss: 0.0374 -
 val_loss: 0.3350
 Epoch 12/20
 363/363 1s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5704 - dense_29_loss: 0.2929 -
 dense_30_RootMeanSquaredError: 0.6089 - dense_30_loss: 0.0371 - loss: 0.3300 -
 val_dense_29_RootMeanSquaredError: 0.5857 - val_dense_29_loss: 0.3087 -
 val_dense_30_RootMeanSquaredError: 0.6805 - val_dense_30_loss: 0.0463 -
 val_loss: 0.3551
 Epoch 13/20
 363/363 1s 1ms/step -
 dense_29_RootMeanSquaredError: 0.5683 - dense_29_loss: 0.2908 -
 dense_30_RootMeanSquaredError: 0.6068 - dense_30_loss: 0.0368 - loss: 0.3276 -
 val_dense_29_RootMeanSquaredError: 0.5485 - val_dense_29_loss: 0.2707 -
 val_dense_30_RootMeanSquaredError: 0.6166 - val_dense_30_loss: 0.0380 -

```

val_loss: 0.3088
Epoch 14/20
363/363          1s 1ms/step -
dense_29_RootMeanSquaredError: 0.5673 - dense_29_loss: 0.2898 -
dense_30_RootMeanSquaredError: 0.6040 - dense_30_loss: 0.0365 - loss: 0.3262 -
val_dense_29_RootMeanSquaredError: 0.5585 - val_dense_29_loss: 0.2807 -
val_dense_30_RootMeanSquaredError: 0.6493 - val_dense_30_loss: 0.0421 -
val_loss: 0.3229
Epoch 15/20
363/363          1s 1ms/step -
dense_29_RootMeanSquaredError: 0.5661 - dense_29_loss: 0.2885 -
dense_30_RootMeanSquaredError: 0.6023 - dense_30_loss: 0.0363 - loss: 0.3247 -
val_dense_29_RootMeanSquaredError: 0.5467 - val_dense_29_loss: 0.2690 -
val_dense_30_RootMeanSquaredError: 0.6175 - val_dense_30_loss: 0.0381 -
val_loss: 0.3072
Epoch 16/20
363/363          0s 1ms/step -
dense_29_RootMeanSquaredError: 0.5643 - dense_29_loss: 0.2867 -
dense_30_RootMeanSquaredError: 0.6007 - dense_30_loss: 0.0361 - loss: 0.3228 -
val_dense_29_RootMeanSquaredError: 0.5477 - val_dense_29_loss: 0.2699 -
val_dense_30_RootMeanSquaredError: 0.6214 - val_dense_30_loss: 0.0386 -
val_loss: 0.3086
Epoch 17/20
363/363          0s 1ms/step -
dense_29_RootMeanSquaredError: 0.5629 - dense_29_loss: 0.2852 -
dense_30_RootMeanSquaredError: 0.5984 - dense_30_loss: 0.0358 - loss: 0.3210 -
val_dense_29_RootMeanSquaredError: 0.5455 - val_dense_29_loss: 0.2678 -
val_dense_30_RootMeanSquaredError: 0.6209 - val_dense_30_loss: 0.0386 -
val_loss: 0.3064
Epoch 18/20
363/363          1s 2ms/step -
dense_29_RootMeanSquaredError: 0.5617 - dense_29_loss: 0.2840 -
dense_30_RootMeanSquaredError: 0.5972 - dense_30_loss: 0.0357 - loss: 0.3197 -
val_dense_29_RootMeanSquaredError: 0.5472 - val_dense_29_loss: 0.2694 -
val_dense_30_RootMeanSquaredError: 0.6014 - val_dense_30_loss: 0.0362 -
val_loss: 0.3056
Epoch 19/20
363/363          1s 1ms/step -
dense_29_RootMeanSquaredError: 0.5603 - dense_29_loss: 0.2826 -
dense_30_RootMeanSquaredError: 0.5961 - dense_30_loss: 0.0355 - loss: 0.3182 -
val_dense_29_RootMeanSquaredError: 0.5542 - val_dense_29_loss: 0.2764 -
val_dense_30_RootMeanSquaredError: 0.6508 - val_dense_30_loss: 0.0424 -
val_loss: 0.3188
Epoch 20/20
363/363          1s 1ms/step -
dense_29_RootMeanSquaredError: 0.5593 - dense_29_loss: 0.2816 -
dense_30_RootMeanSquaredError: 0.5947 - dense_30_loss: 0.0354 - loss: 0.3170 -
val_dense_29_RootMeanSquaredError: 0.6114 - val_dense_29_loss: 0.3363 -

```

```
val_dense_30_RootMeanSquaredError: 0.5949 - val_dense_30_loss: 0.0354 -  
val_loss: 0.3718
```

Evaluating the model performance against testing dataset

```
[76]: eval_results = model.evaluate((X_test_wide, X_test_deep), (y_test, y_test))  
  
weighted_sum_of_losses, main_loss, aux_loss, main_rmse, aux_rmse = eval_results
```

```
162/162          0s 935us/step -  
dense_29_RootMeanSquaredError: 0.5648 - dense_29_loss: 0.2871 -  
dense_30_RootMeanSquaredError: 0.6071 - dense_30_loss: 0.0369 - loss: 0.3240
```

```
[77]: print(weighted_sum_of_losses, main_loss, aux_loss, main_rmse, aux_rmse)  
  
0.32088109850883484 0.28391894698143005 0.036691226065158844 0.5619576573371887  
0.6055104732513428
```

Predicting the results using model

```
[78]: y_pred_main, y_pred_aux = model.predict((X_new_wide, X_new_deep))
```

```
1/1          0s 65ms/step
```

```
[79]: y_pred_tuple = model.predict((X_new_wide, X_new_deep))  
  
y_pred = dict(zip(model.output_names, y_pred_tuple))
```

```
1/1          0s 21ms/step
```

```
[80]: y_pred
```

```
[80]: {'dense_29': array([[0.49514577],  
                        [1.2503316 ],  
                        [3.6080847 ]], dtype=float32),  
      'dense_30': array([[0.4965441],  
                        [1.1439607],  
                        [3.5039604]], dtype=float32)}
```